

ATULA TEJASWI NEERKAJE

Email ◊ Google Scholar ◊ GitHub ◊ LinkedIn ◊ Website

EDUCATION

University of Texas at Austin Ph.D in Computer Science	Austin, TX Jun 2025 -
University of Texas at Austin M.S. in Computer Science, CGPA: 3.86/4.00	Austin, TX Aug 2023 - May 2025
Manipal Institute of Technology Bachelor of Technology (B.Tech) in Computer Science and Engineering, CGPA: 9.49/10	Manipal, India June 2019 - July 2023

EXPERIENCE

Collaborator OpenThoughts-Agent Collaborating with the OpenThoughts-Agents team led by <u>Prof. Ludwig Schmidt</u> to create training data for small, competitive models focused on effective command-line terminal usage.	Oct 2025 - Present <i>Remote</i>
Research Assistant UT Austin Advised by <u>Prof. Sujay Sanghavi</u> . Working on test-time steerability for diffusion language models.	Aug 2025 - Present <i>Austin, TX</i>
Teaching Assistant UT Austin Designing new modules for Data Science Principles/Data Science Lab taught by <u>Prof. Sujay Sanghavi</u> and <u>Prof. Constantine Caramanis</u> .	Aug 2024 - May 2025 <i>Austin, TX</i>
Research Assistant UT Austin Advised by <u>Prof. Eunsol Choi</u> . Worked on multilingual adaptation of Large Language Models and Retriever/Embedding Models.	Jan 2024 - Aug 2024 <i>Austin, TX</i>
Research Intern Indian Institute of Science Advised by <u>Prof. Aditya Gopalan</u> on machine learning and optimization for Large Language Model based problems as part of my Bachelor's Thesis. Worked on efficient fine tuning of language models. Explored active learning strategies for efficient sample selection – showed 1.5x faster convergence on goal-oriented conversation tasks.	Jan 2023 - Jun 2023 <i>Bangalore, India</i>
Research Intern Language, Reasoning and Education Lab, ETH Zurich Advised by <u>Prof. Mrinmaya Sachan</u> on neuro-symbolic methods for Natural Language Processing. Worked on mathematical reasoning problems in Large Language Models. Implemented an In-Context Prompting pipeline inspired from case-based reasoning concepts.	Dec 2022 - Jul 2023 <i>Switzerland (Remote)</i>
Research Intern Université du Québec à Montréal Advised by <u>Prof. Marie-Jean Meurs</u> on an AI-based system for anonymization in corpora of personal data, such as in the legal domain and in the mental health domain. Explored adversarial training and differential privacy for obfuscation of implicit indicators of private information and authorship attributes, in order to release readable, anonymized versions of text corpora for downstream tasks. Also explored Named Entity Recognition and other pseudo-anonymization strategies as part of the pipeline.	May 2022 - Aug 2022 <i>Montréal, Canada</i>

Researcher

Conversational AI and Social Analytics (CAISA) Lab, University of Marburg

Jan 2021 - Jul 2022

Germany (Remote)

Working with **Prof. Lucie Flek** and **Ramit Sawhney** on user-contextual modeling and privacy preservation, mainly for NLP and Computational Social Science problems. Implemented neural network architectures which included Transformers, Graph Neural Networks, and Hyperbolic Neural Networks for these tasks. Studied the impact of differential privacy (DP-SGD) on model fairness, robustness to distribution shifts, and on user-contextual modeling.

AI Division Member

Project MANAS

Feb 2020 - Aug 2022

Manipal, India

The official AI and Robotics team of Manipal, with a team primarily dedicated to building an autonomous car for Indian roads. Worked on Deep Learning and Computer Vision pipelines such as lane detection.

PUBLICATIONS

* indicates equal contribution.

- **Atula Tejaswi***, Litu Rout*, Constantine Caramanis, Sanjay Shakkottai, Sujay Sanghavi. *EntRGi: Entropy Aware Reward Guidance for Diffusion Language Models*. Preprint 2026. [Paper] [Project Page] [Code (Coming Soon)]
- **Atula Tejaswi**, Yoonsang Lee, Sujay Sanghavi*, Eunsol Choi*. *RARe: Retrieval Augmented Retrieval with In-Context Examples*. COLM 2025. [Paper] [Code]
- Vijay Lingam*, **Atula Tejaswi***, Aditya Vavre*, Aneesh Shetty*, Gautham Krishna Gudur*, Joydeep Ghosh, Alex Dimakis, Eunsol Choi, Aleksandar Bojchevski, Sujay Sanghavi. *SVFT: Parameter-Efficient Fine-Tuning with Singular Vectors*. NeurIPS 2024. [Paper] [Code]
- **Atula Tejaswi***, Nilesh Gupta*, Eunsol Choi. *Exploring Design Choices for Building Language-Specific LLMs*. EMNLP 2024 (Findings). [Paper] [Code]
- Victor Petré Bach Hansen*, **Atula Tejaswi***, Ramit Sawhney, Lucie Flek, Anders Søgaard. *The Impact of Differential Privacy on Group Disparity Mitigation*. NAACL 2024 (Findings). [Paper] [Code]
- Ritesh Soun, **Atula Tejaswi**, Ramit Sawhney, Nikolaos Aletras, Preslav Nakov. *RISE: Robust Early-exiting Internal Classifiers for Suicide Risk Evaluation*. LREC COLING 2024. [Paper]
- Samyak Jain*, Parth Chhabra*, **Atula Tejaswi***, Puneet Mathur, Ramit Sawhney, Shivam Agarwal, Preslav Nakov, Sudheer Chava, Dinesh Manocha. *Saliency-Aware Interpolative Augmentation for Multimodal Financial Prediction*. LREC COLING 2024. [Paper]
- Sunny Sanyal*, **Atula Tejaswi***, Jean Kaddour, Abhishek Kumar, Sujay Sanghavi. *Early Weight Averaging meets High Learning Rates for LLM Pre-training*. COLM 2024. [Paper]
- Parth Chhabra*, **Atula Tejaswi***, Shivam Agarwal, Ramit Sawhney, Megh Thakkar, Preslav Nakov, Sudheer Chava. *Learning Through Interpolative Augmentation of Dynamic Curvature Spaces*. SIGIR 2023. [Paper]
- Ramit Sawhney*, **Atula Tejaswi***, Ivan Habernal, Lucie Flek. *How Much User Context Do We Need? Privacy by Design in Mental Health NLP Applications*. ICWSM 2023. [Paper]
- Ramit Sawhney, Megh Thakkar, Ritesh Soun, **Atula Tejaswi**, Vasu Sharma, Dipanwita Guhathakurta, Sudheer Chava. *Tweet Based Reach Aware Temporal Attention Network for NFT Valuation*. EMNLP 2022 (Findings). [Paper] [Code]
- Puneet Mathur, **Atula Tejaswi**, Malika Chhibber, Ramit Sawhney, Fu-Ming Guo, Franck Dernoncourt, Sanghamitra Dutta, Dinesh Manocha. *MONOPOLY: Financial Prediction from MONetary POLicy Conference Videos Using Multimodal Cues*. ACM Multimedia 2022. [Paper] [Code]
- Ramit Sawhney*, Shivam Agarwal*, **Atula Tejaswi***, Nikolaos Aletras, Preslav Nakov, Lucie Flek. *Towards Suicide Ideation Detection Through Online Conversational Context*. SIGIR 2022. [Paper] [Code]
- Ramit Sawhney*, Shivam Agarwal*, **Atula Tejaswi***, Kapil Pathak*. *Orthogonal Multi-Manifold Enriching of Directed Networks*. AISTATS 2022. [Paper] [Code]

- Ramit Sawhney*, **Atula Tejaswi***, Manas Gaur. *A Risk-Averse Mechanism for Suicidality Assessment on Social Media.* ACL 2022. [Paper]
- Ramit Sawhney, **Atula Tejaswi**. *Intermix: An Interference-Based Data Augmentation and Regularization Technique for Automatic Deep Sound Classification.* ICASSP 2022. [Paper] [Code]

AWARDS

CLS Fellowship 2025

Selected for a fellowship awarded by **Max-Planck ETH Center** to pursue a PhD at ETH Zurich. Declined to pursue PhD at UT Austin.

NeurIPS Travel Award 2024

Awarded a travel grant by NeurIPS 2024 to present my paper SVFT in Vancouver, Canada.

Google Research Week 2023

Selected to be a part of the ML Foundations track in a three day research week comprising of lectures and discussions with faculty and researchers from Google, DeepMind, ETH Zurich, etc.

Mitacs GRI 2022

Awarded a fully-funded scholarship by Mitacs Inc., Canada to pursue a summer research project at the Université du Québec à Montréal on anonymization and bias reduction in AI & Natural Language systems.

DAAD WISE 2022

Selected for the DAAD WISE research scholarship award for Summer 2022 (Declined).

Mahindra Rise Prize Challenge

Project MANAS won the Million Dollar Mahindra Rise Prize Challenge, placed top 13 out of 153 teams in India.

GRADUATE COURSEWORK

Convex Optimization (CS395T), Advances in Deep Generative Models (CS395T), Reinforcement Learning - Theory and Practice (CS394R), Generative Models in Machine Learning (ECE381V), Topics in Natural Language Processing (CS395T), Database Systems (CS386D), Cryptography (CS388H).

ACTIVITIES

Reviewer NeurIPS 2024-2025, ICML 2025, ICLR 2025-2026, ARR, AISTATS 2022
Volunteer AISTATS 2022